


# Project 1

Emil Engelstad Moghaddam

October 10, 2020

## Abstract

In this project I have studied how data from the Frankfunction( $x, y$ ) can be approximated with the least squares method with polynomials of  $x$  and  $y$  as predictive variables. When training the model on a subset of the data, the predictive ability of the model, measured in the  $R^2$  score and mean squared error keeps improving as we add further polynomials of higher degrees to the design matrix. However the models ability to predict the rest of the data, the test data, peaks when the contains polynomials of  $x$  and  $y$  of maximum 4 to 6. We decompose the MSE into bias and variance to study why this happens, and see that the variance starts increasing in this region. We add the penalty terms characteristic of ridge and lasso regression before calculating the models, and see that the variance stays low as polynomial degrees increase, which result in the MSE staying consistently low. We apply the three methods of regression to terrain data, using polynomials of maximum degree 5 in the design matrix. Normal linear regression and ridge regression is able to replicate the terrain data with an  $R^2$  score of .56, while lasso regression replicates with an  $R^2$  score of .54.



# 1 Introduction

In this project we are given points  $x$  and  $y$  and a outcome variable  $Y$ . In the first part of the project the outcome variable  $Y$  is the value of the Frank function of  $x$  and  $y$ , in the second part the value is the height of the point  $(x,y)$  in the terrain data. What we are interested in is finding a formula that transforms the points  $(x,y)$  as close as possible to the outcome variable  $Y$ . We designate the vector of the points calculated by  $\tilde{y}$ . We call the numerical distances between the outcome data and the predicted value of the model the residuals.

When we use the least squares method we put the variables we want to use in a design matrix  $X$ . In this case it is powers of  $x$  and  $y$ . We call  $\beta_j$  the amount we use of variable  $j$  used to approximate the data, and designate  $\beta$  as the vector of all the beta values. We can now derive an analytical expression for the squared values of the residuals which we call the cost/loss function. The smaller the residuals the better our functions, therefore we want to find optimal beta values to minimize this function. Mathematically we can write it as:

$$Costfunction = |Y - X\beta|^2 = [Y - X\beta]^T [Y - X\beta] = \sum (y_i - \sum x_{ij}\beta_j)^2 \quad (1)$$

When analyzing data it is normal to make the assumption that the data follows some sort of a distribution dependent on some explanatory variables and that the data contains some noise. When an algorithm is overly sensitive to the data and isn't able to extract the distribution from the noise we say that the algorithm is overfitting. We can get an indication that an algorithm is overfitting by either seeing that there is a big discrepancy between the results on training data versus test data, or by noticing that the model changes a lot when we train it on different samples of our data. One solution to this problem is by penalizing higher beta values. Ridge and lasso regression are two different implementations of this solution with their own loss/cost functions:

$$Costfunction^{Ridge} = |Y - X\beta|^2 + \lambda|\beta|^2 = \sum (y - \sum x\beta)^2 + \lambda \sum \beta^2 \quad (2)$$

$$Costfunction^{Lasso} = |Y - X\beta|^2 + \lambda|\beta| = \sum (y - \sum x\beta)^2 + \lambda \sum |\beta| \quad (3)$$

$\lambda$  is the amount we choose to penalize the beta values with, and we can change this value to optimize our results.

## 2 Theoretical model and technicalities

In this analysis I have used two types of measures to quantify how well our models predict the data; the mean squared error and R - Squared. In both cases we start by taking each residual, the distances from what our model predicts to

the actual value of our data at a specific point, and squaring the value. With the mean squared error we divide by number of points, which lets us datasets of different sizes.

$$MSE = \frac{(y - \tilde{y})^2}{n} \quad (4)$$

However MSE doesn't provide any intuitive understanding of how good the model is. With the R-squared we instead divide by the total sum of squares, which is proportional to the variance, and is the squared distance of each point from the mean. The fraction tells us how much of variance of the data isn't explained by the model. The mathematical formula below shows how to turn this fraction into a number between one and zero which tells us how much of the variance is explained by the model.

$$R^2 = 1 - \frac{(y - \tilde{y})^2}{(y - \mu)^2} \quad (5)$$

## 2.1 Pseudo code

The general form of the project follows the form of the pseudo code 1 page below. We choose what type of linear regression model we want to use, and what form of sample methods we want to use. Then we calculate and evaluate models for different design matrices.

## 2.2 Least Squares theory

Least squares is a method of where we multiply each of our  $p$  independent variables with values,  $\beta_0$  to  $\beta_{p-1}$ , to calculate  $\tilde{y}$ . The method tries to minimize the sum of all the residuals. The first concise exposition of the method of least squares was published by Adrien-Marie Legendre in 1805. The method is however usually credited to Carl Fredrich Gauss. The formula for the optimal beta values, that minimize the squared residuals, can be calculated by setting the gradient of the residuals for each of the beta values equal to zero. Another clever way to find the optimal beta values is by realizing that the residuals are in the null space of the transpose of the matrix containing all the independent variables, making the equality:

$$\begin{aligned} X^T(Y - X\beta) &= 0 \\ X^T X\beta &= X^T Y \\ \beta &= (X^T X)^{-1} X^T Y \end{aligned}$$

The projection matrix is given by  $X(X^T X)^{-1} X^T$  and  $\tilde{y}$  by  $X\beta$ . In the project the beta values have been calculated with the training data and  $\tilde{y}$  with both training and test data. The last part of the analysis uses the singular value decomposition version of least squares.

The singular value decomposition theorem states that any  $n \times p$  matrix  $A$  with  $p \leq n$  has some orthonormal vectors  $v_i$  in  $V$ , that  $A$  transforms into another

---

```

 $\vec{x} \leftarrow$  random numbers in range[0,1]
 $\vec{y} \leftarrow$  random numbers in range[0,1]
 $\vec{data} \leftarrow \text{Frankfunction}(\vec{x}, \vec{y})$ 

 $X \leftarrow [\vec{1}, \vec{x}, \vec{y}, \vec{x} \times \vec{y}]$  design matrix

for  $p \leq$  Maximum polynomial do
     $X \leftarrow +\vec{x}^p, \vec{y}^p$ 

    if Bootstrap method then

         $\text{TrainTestSplit}(X, \text{Data})$ 
        for  $b \leq$  nr of Bootstraps do

             $X_{\text{bootstrap}}, \text{Data}_{\text{bootstrap}} \leftarrow \text{bootstrap}(X_{\text{training}}, \text{data}_{\text{training}})$ 
             $\text{model}_{p,b} = \text{Normal/Ridge/Lasso}(X_{\text{bootstrap}}, \text{Data}_{\text{bootstrap}})$ 
             $R^2_{\text{Training/test}} = R^2(\text{Data}_{\text{training/test}}, \text{model}_{p,b})$ 
             $\text{MSE}_{\text{Training/test}} = \text{MSE}(\text{Data}_{\text{training/test}}, \text{model}_{p,b})$ 

        end for
    end if

    if K - fold regresstion then

         $\text{SplittK}(X, \text{Data})$  Splitting into k equal parts
        for  $k \leq K$  do

             $\text{Model}_k = \text{Normal/Ridge/Lasso}(X_{\neq k}, \text{Data}_{\neq k})$ 
             $R^2_{\text{test/Training}} = R^2(\text{Data}_{k/\neq k}, \text{model}_k)$ 
             $\text{MSE}_{\text{test/Training}} = \text{MSE}(\text{Data}_{k/\neq k}, \text{model}_k)$ 

        end for
    end if
end for

 $\text{Bias}_p = \text{MSE}(\text{data}, \text{mean}(\text{model}_p))$ 
 $\text{Variance}_p = \text{var}(\text{models}[p])$ 
 $\text{plot}(R^2_{\text{training}}, R^2_{\text{test}}, \text{MSE}_{\text{training}}, \text{MSE}_{\text{test}})$ 

```

---

set of orthonormal vector U, stretched in proportion to the diagonal matrix  $\Sigma$ .

$$\begin{aligned} AV &= U\Sigma \\ A &= U\Sigma V^t \end{aligned}$$

We define the matrix D as the values of  $\Sigma$  not equal to zero.

$$\begin{aligned} \beta &= (X^T X)^{-1} X^T y \\ \beta &= ((UDV^T)^T (UDV^T))^{-1} (UDV^T)^T y \\ \beta &= (VDU^T UDV^T)^{-1} (VDU^T)^T y \\ \beta &= VD^{-2} V^T VDU^T y \\ \beta &= VD^{-1} U^T y \end{aligned}$$

### 2.3 Confidence intervals

To calculate confidence intervals for the beta values we need to know how they are distributed. If we assume that  $\varepsilon_i \sim N(0, \sigma^2 I)$ , we can derive a normal distribution for the beta values. Note that  $\sigma^2 I$  implies uncorrelated error terms.

$$\hat{\beta} \sim N(E[\hat{\beta}], Var[\hat{\beta}])$$

$$\begin{aligned} \hat{\beta} &= (X^T X)^{-1} X^T Y = (X^T X)^{-1} X^T (X\beta + \varepsilon) \\ &= \beta + (X^T X)^{-1} X^T \varepsilon \end{aligned}$$

$$E[\varepsilon] = 0 \text{ and } E[\beta] = \beta \text{ therefore } E[\hat{\beta}] = \beta.$$

$$\begin{aligned} Var[\hat{\beta}] &= Var[(X^T X)^{-1} X^T \varepsilon] = (X^T X)^{-1} X^T \varepsilon ((X^T X)^{-1} X^T \varepsilon)^T \\ &= \sigma^2 (X^T X)^{-1} X^T ((X^T X)^{-1} X^T)^T = \sigma^2 (X^T X)^{-1} \end{aligned}$$

From this normally distributed Beta values we can easily calculate confidence intervals and we can see that they are proportional to the variance of the variables  $X^T X$  and the variance of the dataset. Note; i have not shown that the betas are normally distributed, only the variance and the expected value. I have the sigma of the noise added to the frank function to when calculating the confidence intervals in project 1.

### 2.4 Ridge regression

In ridge regression we want to minimize the sum of residuals squared plus a  $\lambda$  times the sum of betas squared. We penalize the beta values to prevent overfitting. We generally don't want to penalize beta 0. To avoid this we set  $\beta_0$  equal to  $\bar{Y}$  and define a new variable  $Y^* = Y - \bar{Y}$ . Then we calculate the rest of the beta values with  $Y^*$  as the dependent variable.

Trying to minimize the sum of square residuals plus the  $\lambda$  times sum of squared beta values equals is equivalent to minimizing the sum of squared

residuals while only picking beta values from a p dimensional ball.  $\sum \beta^2 \leq t$   
The ball is placed in a p dimensional grid with each axis representing the size  
one beta value. Imagining the optimal beta value is somewhere outside that  
ball, the growing p dimensional ellipse from optimal value to the constraint  
ball will not hit one of the axes first unless unlikely case the optimal beta  
value already was zero. Therefore beta values will never shrunk to zero.

We can show that the second derivation of the loss fuction with respect to  
beta is positive definite. This tells us that the analytic expression for the beta  
gives a true minimum for the loss function, as a contrast to a saddle point.

Note i have divided by 2 to make the math cleaner.

$$\frac{\delta L}{\delta \beta} = \beta(X^T X) + \beta \lambda I - X^T y$$

$$\frac{\delta L}{\delta \beta \delta \beta^T} = (X^T X) + \lambda I$$

The covariace matrix is positive semidifinite, and acoridng to (Lemma 14.2.4 of  
Harville, 2008), adding a positive lambda makes it positive definite. This is  
also what we use for pseudo inverse when the  $X^T X$  is not invertible.  
In the last task of the project we use the othodinal decomposition version of  
rige regression

$$\beta^{Ridge} = (X^T X + \lambda I)^{-1} X^T y$$

$$\begin{pmatrix} (VD^2V^T + \lambda I)^{-1} & VD^2U^T y \\ (VD^2V^T + \lambda V^T V)^{-1} & VD^2U^T y \\ (V(D^2 + \lambda)V^T)^{-1} & VD^2U^T y \\ V(D^2 + \lambda)^{-1}V^T & VD^2U^T y \end{pmatrix}$$

$$V \frac{D}{(D^2 + \lambda)} U^T y.$$

Notice that we know from the definition given earlier of singular value  
decomposition how X extends and rotates every vector in V. The result is that  
the shrinkage factor becomes  $\frac{D^2}{(D^2 + \lambda)}$ . The further the vector is extended the  
less it is reduced. This shows ut that it is important that the vectors are  
scaled for ridge to work properly.

## Lasso regression

Lasso regresstion uses L1 regularization. The additional penalty term is  $\lambda$  times  
the sum of the absolute value of each of the betas. The optimal beta values for  
lasso can not be soled for analytically. The advantage of lasso is that it does  
feature selection by setting beta values to zero. The penalty of lasso is  
equivalent with only picking from a p dimensional square. Here the contours  
representing equal increase in error from the minimal sum of squared residuals  
may hitt a corner of the square.

## Bias-variance tradeoff

$$MSE = E[(data - model)^2]$$




Figure 1: Here we see the difference of lasso and ridge in two dimensions

$$\frac{E[data^2] - 2 \times E[data \times model] + E[model^2]}{E[data^2] + E[model^2] - 2 \times E[data \times model]}$$

Here we use the classical result from statistics that  $Var(X) = E[X^2] - E[X]^2$ .

Moving the last term and substituting into the right side we get:

$$Var(data) + E[data]^2 + var(model) + E[model]^2 - 2 \times E[data \times model]$$

$$\begin{aligned} &\text{We use: } var(data) = \varepsilon^2 \text{ and } E[data] = data \text{ and simplify} \\ &E[data]^2 - 2 \times E[data \times model] + E[model]^2 = (data - E(model))^2 = bias^2 \\ &\implies \\ &MSE = \varepsilon^2 + bias^2 + variance \end{aligned}$$

What we see here is that the mean squared error is a function of bias and variance. If we don't use the necessary variables to predict the data, like trying to predict house prices without house size, then each model we calculate with different sets of data will predict badly. So will the mean of these models and the bias will be high.

What is worth noticing is that even though the mean model might predict the data well, the mean squared error might still be high. This can happen when the model is too flexible, using too many unrelated variables. Unrelated variables can be used to further reduce the error of a model on a dataset, but when that model is applied to another dataset it will do worse than the original dataset without the unnecessary terms. Generally we want both low bias and low variance.

## Resampling techniques

In this project i have implemented the resampling techniques; the bootstrap and k-fold regression.

K-fold regresstion works by splitting the data into k parts. Each part is used as test data exactly once, the k-1 other parts are used as training data for the model. The advantage with k-fold regression is that we get to use the entire dataset.

With bootstrap we splitt data into test and training data. Instead of using the trainingdata directly, bootstrap makes n resamples of the data. Each datapoint has equal chance of being picked, and picking a datapoint does not reduce its chance of being picked again.

The advantages of bootstrap is that one can calculate confidence intervalls, and check the stability. According to the original developer of the bootstrap 50 resamples should lead to good estimate.

Note: Randomization of datapoints for k-fold has only been implemented in the last task as the values are already randomly generated in the first parts.



### 3 Results and discussion

#### Normal linear regression




Figure 2:  $R^2$




Figure 3: Mean squared error

On the graphs we see that the  $r^2$  score and mean squared error of the training data keeps improving for the test data as polynomials of higher degrees are added to the predictive variables. We also see that this is not the case for the test data. Here the model keeps improving until we reach a polynomial of degree 4, and then quality of the predictions decreases exponentially. At the peak the model has a  $r^2$  score of .93. We also see that the model predicts training data more accurately than the test data for all degrees of polynomials.

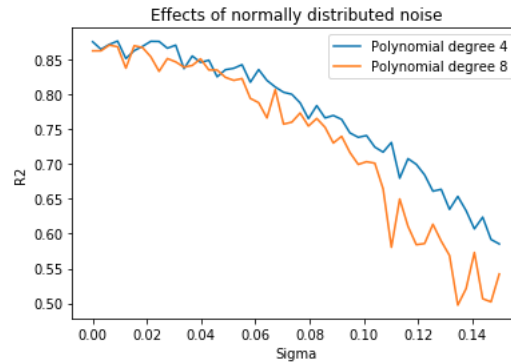


Figure 4: Noise

The graph shows the effect of normally distributed noise. The graph shows that a model that uses more variables might react slightly more to the noise of the data.




Figure 5: Confidence interval for  $\beta$

This figure shows a 99 percent confidence interval for the beta values when using polynomials of  $x$  and  $y$  up to degree 3 as predictive variables. The results are from a analytic calculation using  $(X^T X)^{-1} \times \sigma^2$ . The beta values have been calculated with 100 random resamples of the data, and the results look the same.

b




Figure 6: df




Figure 7: df

The same pattern is present when analysing the data with 10 fold normal linear regression. The test MSE and R2 score appears to be more slitley stable than in the previous case.

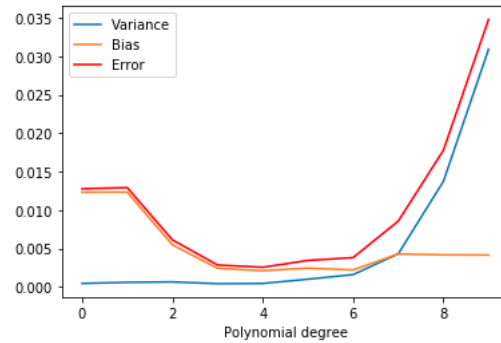


Figure 8: This is an image from a text that uses color to teach music.

The graf illustrates the bias - variance tradeoff. Notice that the error term is the mathematical sum of the bias term and the variance. On the model we can se that the bias term decreases as we add more and more predictive variables of higher polynomial degress to the model. The bias represents the mean squared error of the avarage of 200 bootstrap created models. The variance is low for lower polynomials and increases exponenally as further polynomials are added. The variance describes the deviance of the models from the mean model. We se that for the low variace error closly follows the bias. For higher polynomials even though the mean model predicts well, the induvidual models deviate far from this model, and they are oversensitive to the data and they predict the the test data badly.

## Lasso and Ridge




Figure 9:  $R^2$

The results are calculated using a 10 fold split. Both functions perform well even when using many unnecessary higher polynomials as variables in the design matrix. Here we see that the ridge regression predicts the franke function slightly better than the lasso regression. The reason the functions stay stable at higher polynomials is present in the bias variance tradeoff.

## Bias variance




Figure 10: Ridge




Figure 11: Lasso

These graphs show the bias-variance tradeoff for ridge and lasso regression.

The graph shows that the bias, the prediction error of the mean model, improves until we add polynomials of degree 6. This is true for both models.

From there on the bias term is stable. This is the same as in normal linear

regression. The variance is low and stable for both models through all degrees of polynomials, as a contrast to normal linear regression where the variance increases exponentially with added polynomials. The variance is slightly higher for polynomials 10-14 with ridge regression, but the bias and error term is lower.

$\lambda$




Figure 12:  $\text{Ridge}(\lambda)R^2$




Figure 13:  $\text{Lasso}(\lambda) \text{ MSE}$

The graphs show how the results of both lasso and ridge depend upon a weight  $\lambda(\text{lambda})$ . The ability of the ridge and lasso model to predict the test data keeps improving when increasing lambda, until a point where they drastically decrease. On the right graph we are using the  $R^2$  metric and therefore interested in finding the top of the graph. On the left graph we are using the mean squared error and therefore interested in finding the bottom of the graph which indicates the optimal lambda value.

## Terrain data




Figure 14: This is an image of how three methods reproduce the terrain data.

Here the picture has been recreated using all three methods. With zero noise added to the terrain data normal linear regression and ridge regression replicate the results with an  $R^2$  score of .56. Here there has been added a significant

amount of normally distributed noise and the  $R^2$  score is .49. Lasso does slightly worse both with and without noise.

## 4 Conclusions

Ridge seems to overall to perform the best of the three models when wanting have a model that predict well and one is using some variables uncorrelated with the data one, and one is not interested in variable selection. Both lasso and Ridge does a good job of penalizing reducing the variance of the models for the higher polynomials

## 5 Appendix with extra material

Harville, D. A. (2008). Matrix Algebra From a Statisticians Perspective. Springer, New York.

Hastie, T., Friedman, J., and Tibshirani, R. (2009). The Elements of Statistical Learning. Springer

Stigler, Stephen M. (1981). "Gauss and the Invention of Least Squares"

Bretscher, Otto (1995). Linear Algebra With Applications (3rd ed.)